

# Modeling Global and local Codon Bias with Deep Language Models

M. Stanley Fujimoto\*, Paul M. Bodily, Cole A. Lyman, J. Andrew Jacobsen, Quinn Snell, Mark J. Clement  
 Computer Science Department  
 Brigham Young University  
 Provo, Utah 84604  
 \*sfujimoto@gmail.com

**Abstract**—Codon bias, the usage patterns of synonymous codons for encoding a protein sequence as nucleotides, is a biological phenomenon that is not fully understood. Several methods exist to represent the codon bias of an organism: codon adaptation index (CAI) [1], individual codon usage (ICU), hidden stop codons (HSC) [2] and codon context (CC) [3]. These methods are often employed in the optimization of heterologous gene expression to increase the accuracy and rate of translation. They, however, have many shortcomings as they don't take into account the local and global context of a gene. We present a method for modeling global and local codon bias through deep language models that is more robust than current methods by providing more contextual information and long-range dependencies.

**Keywords**—codon usage bias, deep learning, language modeling, machine learning

## I. INTRODUCTION

Codon usage bias describes the usage pattern of synonymous codons, different codons that encode the same amino acids, in an organism. Studies have observed that differences exist in frequencies of synonymous codon usage across different organisms [4], [5], [6]. The genome hypothesis of codon bias is that different organisms have a distinct codon bias from other organisms [5]. Additionally, the codon bias within a particular organism is complex as even within the same genome codon usage varies amongst genes [7], [5].

Codon usage influences many aspects of biology: RNA secondary structure [8], gene expression [9], speed of translation elongation and protein folding [10]. Thus, codon usage bias is an important phenomenon to understand. Specifically, in synthetic biology, codon usage bias is critical in optimizing heterologous gene expression. A protein of interest is usually presented as an amino acid sequence. The protein must be reverse-translated (backtranslated), converted from an amino acid sequence to a nucleotide sequence. The backtranslated sequence is then injected into a biological vector (e.g. *Escherichia coli*) for heterologous gene express the protein. Engineers seek to use codons that correspond to the amino acids of the protein of interest that maximize the expression of the protein to reduce production time and to use smaller culture volumes. A biologically unviable protein may result if the nucleotide sequence for a protein is transcribed into an RNA transcript that degrades prematurely, does not fold in the correct manner to confer correct function or is expressed in small quantities. Welch et al. found that proper selection of codons can result in 100-fold differences in expression of genes in *E. coli* [11].

The cause of codon bias is unclear. Hershberg and Petrov discuss that there are two general classes of explanations of codon bias: selectionist and mutational [12]. The selectionist explanation states that codons are selected to maintain efficiency and/or accuracy of protein expression. This explanation is supported by the correspondence of *preferred codons*, frequently used codons, and tRNAs that occur more abundantly within an organism [13], [14]. The mutational explanation states that codon bias exists because mutational patterns are nonrandom.

### A. Measuring Codon Bias

Several methods exist to statistically analyze codon bias in different organisms, including: codon adaptation index, frequency of optimal codons and relative codon adaptation index. These methods are used to analyze the codon usage bias to predict gene expression levels. For example, the codon adaptation index (CAI) [1] is the most widespread method for codon usage bias analysis. CAI is a method that measures the deviation of the coding sequence of a gene compared to a reference set of genes. The ability to predict gene expression does not provide a method for predicting which of the synonymous codons should be used when given the amino acid sequence of a gene.

### B. Codon Selection and Bias Modeling

Many algorithms and applications exist that are designed to solve the synonymous codon selection problem for backtranslating a protein sequence into the appropriate coding sequence while optimizing for translational accuracy and synthesis yield for a given bacterial vector.

Many algorithms model codon bias with respect to different facets of biology. Some algorithms attempt to mimic the codon usage found throughout all, or a subset of, genes in an organism. Frequencies for each set of synonymous codons can easily be calculated given the *coding sequence* (CDS) for a genome. Using the most frequently occurring codon of a set of synonymous codons can then be used to backtranslate amino acid sequences. Additionally, probabilistically selecting codons is also an option given usage rates of synonymous codons. Codon selection is probably not determined by a hard rule selecting the most frequently used synonymous codon as evidenced by the usage of other, rare codons. Codon selection is, as proposed by the mutational explanation of codon bias, a non-random event making the probabilistic approach inappropriate. Research has reported that the choice of high-frequency

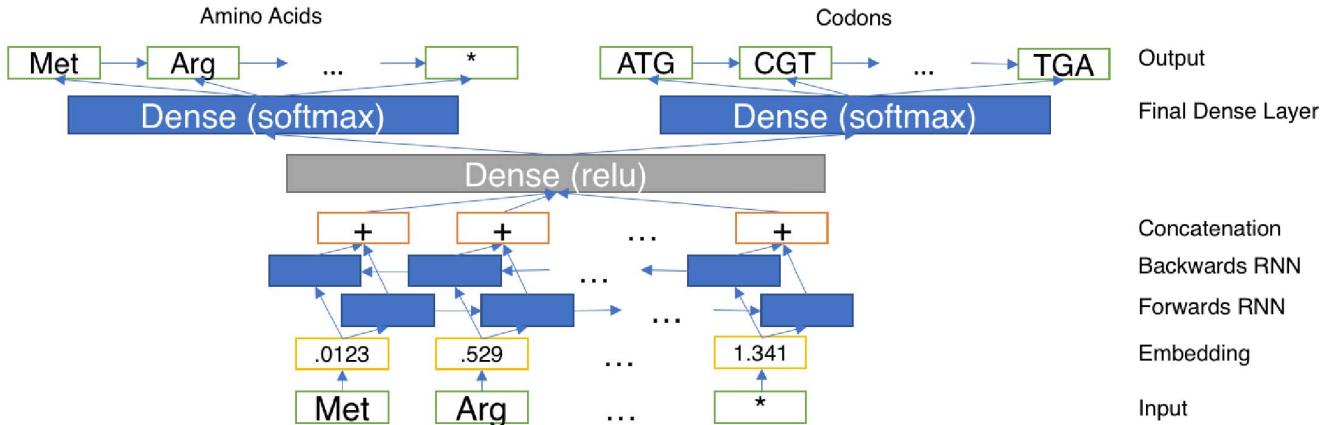


Fig. 1. Basic model architecture. The model is very similar to that used in Neural Machine Translation with the exception of tighter coupling between the encoder and decoder. Our model also features two target outputs: amino acid sequences and codon sequences. The final model that we used varied from this with additional recurrent and classification layers.

codons contributes less than other factors when evaluating translational efficiency [15], [16], [17], [18].

Another method for codon selection is to select codons that correspond to tRNAs that are abundant in the host organism. Intuitively, this method seems appropriate because it would allow for the rapid translation of a protein and gives biological significance to the previously mentioned method. Studies have found, however, that there are circumstances where slowing the rate of translation is desirable. This may be the case because during translation the protein is folding and by slowing down translation the correct structures are able to form. If translation occurs too quickly, the protein may not fold in the correct manner.

Two applications that implement methods for codon selection for backtranslation are DNAWorks and GeneGPS.

1) *DNAWorks*: DNAWorks is a tool used in synthetic gene engineering for backtranslating amino acid sequences given a target vector for gene expression [19]. It is designed to take an amino acid sequence as input and a table that specifies the usage frequency of synonymous codons of the intended biological vector for expression. The usage table is derived from a set of genes that contain properties (e.g. highly expressed) which the user wishes to confer on the backtranslated sequence.

DNAWorks' method for optimizing the codons in the nucleotide sequences relies on choosing amongst a set of synonymous codons the most frequently observed codon as annotated in the input table. Post-processing can modify the sequence further to reduce unwanted structures (hairpins, etc.) from forming during PCR. Their main method, however, of modeling the codon bias of a particular organism is limited as it is represented solely as a table of synonymous codon usage frequencies and contradicts what is now known about codon bias.

2) *GeneGPS*: Another application that exists is GeneGPS by ATUM. GeneGPS is based on the work done by Welch et al. [11] where they found that, in *E. coli*, that the expression levels of proteins were most strongly dependent on codons with corresponding tRNAs that are highly charged during amino

acid starvation and not the most commonly used codons in highly-expressed *E. coli* proteins. Codon selection for high expression should maintain high levels of charged tRNAs and minimize levels of uncharged tRNAs. The results from Welch et al. show that mimicking the host's codon bias or using high CAI does not result in optimal expression levels.

## II. METHODS

We propose a method of modeling codon usage for a particular organism that is more robust than current methods. This method seeks to capture the global and local codon bias of an organism in a richer manner than other methods. The method that we have developed selects codons based on short- and long-range contextual information. This differs from other methods that use global codon usage rates. By using the known coding sequence of existing proteins, our model learns the appropriate synonymous codon for an amino acid while considering the entirety of the protein sequence. Using contextual information allows our model to learn when to use certain codons as selection may differ when forming different types of protein structures or how the nucleotide sequence will influence mRNA secondary structure formation. This contextual information gives a much richer model of codon bias for an organism.

We present a method for modeling the codon bias of an organism by using a *bi-directional recurrent neural network* (BRNN) with either *long short-term memory* (LSTM) or *gated recurrent unit* (GRU) memory cells [20], [21], [22].

### A. Datasets

We trained base models using the 159 different *Escherichia coli* strains. These data were gathered from NCBI using the Entrez API through BioPython [23] which allows for automated querying and retrieval of both *coding sequences* (CDSs) and corresponding translated amino acid sequences. The model was then fine-tuned for 4 of the *E. coli* strains and for the bacteria *Staphylococcus aureus* and *Yersinia enterocolitica* (see Table III for strain NCBI accessions).

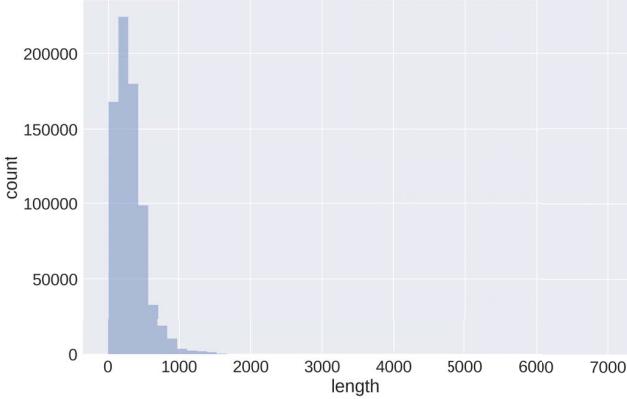


Fig. 2. Histogram of protein sizes used from the 159 pooled *E. coli* genomes. Most proteins used fall under 1,000 amino acids in length with the longest protein 6,926 amino acids and the shortest with 15 amino acids.

TABLE I. SEQUENCE LENGTH FILTERING FOR THE *E. coli* DATASET SHOWING THE NUMBER OF TRAINING INSTANCES AFTER SEQUENCE LENGTH FILTERING HAS BEEN APPLIED.

Seq Len	Train/Validate	Test
100	83531	20883
250	278690	69673
500	516151	129038
1000	586227	146557

### B. Deep Language Model

To model the codon bias of a specific organism we use a method similar to those used in *Natural Language Processing* (NLP) that uses deep neural networks for natural language translation and for language modeling. We treat amino acid sequences as a source language and the corresponding codon sequences as the target language.

Neural machine translation (NMT) is a method used for natural language translation that uses neural networks [24], [25]. The translation task is to take a sentence in one language as input and output the sentence in another language. Neural machine translation relied on an *encoder/decoder* scheme to solve the task. The input sentence is encoded using a neural network, the *encoder*, into a vector representation of float values. A separate neural network, the *decoder*, then takes the encoded values as its input and outputs a sentence in the target language.

### C. Network Architecture

Our basic neural network architecture is a BRNN trained with amino acid sequences as inputs and the known codons as target outputs (see Figure 1). We tried variations of this network architecture by changing the number of recurrent layers and classification layers (see Table II). LSTM or GRU cells are specified at run-time by the user for use in the BRNN. Each separate amino acid in the input sequence is regarded as a time-step in the overall sequence by the network. Our model architecture differs from those used in NMT where NMT does a complete encoding of the input and then passes that complete encoding to the decoder. This results in outputs that can have lengths that differ from the input (as can be the case in natural language translation). For the amino acid

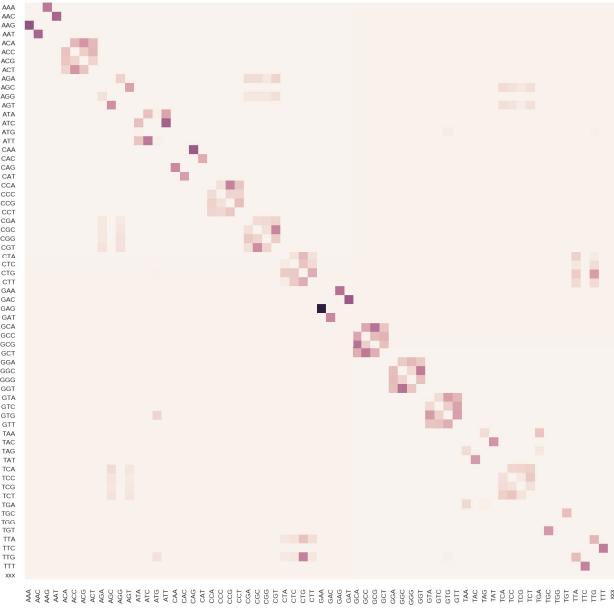


Fig. 3. Heat map of erroneous codon substitutions by the sequence length 100, no dropout model. The x-axis represents the correct codon and the y-axis represents the codon incorrectly substituted for by the model.

detranslation problem, we know that there will be as many codons as there are input amino acids. Thus, we construct our network to produce an output at every time-step instead of encoding the entire sequence first.

Additionally, we include a second task for the network: output the input amino acid. Adding this task pushed the network to act as an autoencoder. This was added to help the network output only codons that encode the correct amino acid and if the incorrect codon is selected it should be a synonym of the correct codon.

### D. Fine-Tuning

After training the base model, we fine-tuned models for specific *E. coli* strains and for other bacteria. Fine-tuning involves taking the base model and continuing training using a more specific dataset [26].

Different variations of fine-tuning exist. One method is to freeze the weights in the recurrent layers and leave the classification layer weights unfrozen and trainable. This results in the final fully-connected layer being modified and the recurrent layers remaining the same as the base model during fine-tuning.

Fine-tuning can also be accomplished by replacing the classification layers of the network and completely retraining them while the other layer weights remain frozen. This may be beneficial if the sequences used for fine-tuning vary greatly from the original dataset used to train the network. This still leverages the learned features from the network but allows it to learn final codon selection in a manner much more specific to the specified sequences.

Fine-tuning allows us to build networks that are specific to a particular species or even for a specific strain of a bacteria.

TABLE II. MODEL PARAMETERS AND TESTING ACCURACY. NOTE THAT THE TRAINING ACCURACY IS USUALLY LOWER THAN THE VALIDATION ACCURACY WHEN DROPOUT IS USED BECAUSE A PORTION OF THE NETWORK IS UNUSED DURING TRAINING WHILE THE FULL NETWORK IS USED ON THE VALIDATION DATASET WHILE.

RNN Cell Type	# RNN Layers	Network Architecture					Training		Validation		Testing	
		Embedding Size	Network Width	Batch Size	Epochs	Dropout	Seq Len	CDS Acc	AAS Acc	CDS Acc	AAS Acc	CDS Acc
LSTM	2	64	512	128	10	0%	100	98.63%	100.00%	97.48%	100.00%	97.35%
LSTM	2	64	512	128	10	0%	250	97.60%	100.00%	96.98%	100.00%	97.04%
LSTM	2	64	512	128	10	0%	500	96.78%	100.00%	96.65%	100.00%	96.70%
LSTM	2	64	512	32	2	0%	1000	94.05%	100.00%	94.47%	100.00%	-
LSTM	2	64	512	128	10	20%	100	97.07%	100.00%	97.31%	100.00%	97.15%
LSTM	2	64	512	128	10	20%	250	94.02%	100.00%	95.51%	100.00%	95.55%
LSTM	2	64	512	96	10	20%	500	88.65%	100.00%	93.79%	100.00%	93.79%
LSTM	2	64	512	32	2	20%	1000	79.86%	100.00%	91.08%	100.00%	-

We perform fine-tuning instead of re-training a new network to reduce training time. The base model is trained with all data from different species/strains pooled. After the base model is trained, we then select a specific species or strain to fine-tune with resulting in a model that can train quicker by leveraging previous work.

### E. Additional Heuristics

While this approach showed high codon selection accuracy, there were still instances where the model would choose codons that encoded the wrong amino acid with the addition of the amino acid target function to the network (see Section II-C). We correct for this by taking the softmax probabilities for the codons, mask probabilities for any codon that does not encode the correct amino acid and select the highest probability from the remaining, valid codons. Amino acid accuracy is guaranteed to be perfectly accurate with this post-processing measure but does not make a significant difference in overall codon selection accuracy. This is most likely because instances where a codon that encodes the incorrect amino acid are rare occurrences.

This heuristic can be applied to all positions in the predicted sequence except the start codon. While methionine is usually the first amino acid, it is not always encoded by AUG. Alternative start codons may be used and still encode a methionine. This is due to a special initiator tRNA that is used during translation [27]. Recently, more alternative start codons have been discovered [28]. Thus, we leave the network with the ability to choose any codon as a start codon as it is unclear which codons are valid start codons as there may be yet undiscovered start codons.

### F. Implementation

Model training was done using NVIDIA K40, Titan X Maxwell and Titan X Pascal GPUs. Our models were built using the Keras [29] neural network library. All our models were developed using Theano [30] as the backend for Keras. The model can be run using different parameters for the type of recurrent cell, dropout in the recurrent cells, the number of weights for different layers and number of recurrent layers (see Table II for tested parameters). All models were trained using the Adam optimizer [31].

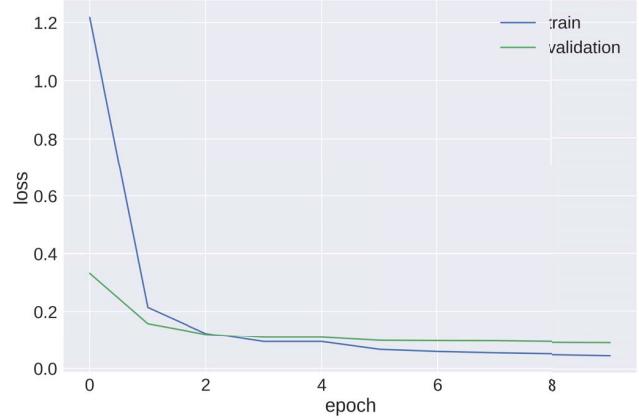


Fig. 4. Loss over epochs for the SEQ LEN 100 dataset. The same loss trend is observed for all training parameters tested.

## III. RESULTS AND DISCUSSION

### A. Base Model

The results of training the base model using the pooled data from 159 *E. coli* genomes can be seen in Table II. Data was pooled to provide enough data to train the model where a single strain did not. In total, the pooled dataset consisted of 742,494 sequences. Characteristics of the dataset can be seen in Figure 2.

A number of different models were trained with different parameters. Selecting for different maximum sequence lengths can greatly affect training time as long sequences take longer to process. We tested different sequence lengths (100, 250, 500 and 1,000) and achieved high accuracy that decreased as sequence length increased. Filtering the dataset based on sequence length resulted in different dataset sizes used for training, validation and test (see Table I). Although the number of instances does not grow dramatically when filtering sequences of 500 and 1,000 amino acids, the overall amount of data increases dramatically.

Reported training accuracy is lower than validation accuracy for models using dropout. This is because during training with dropout, a portion of the network is unused. This results in lower performance. When calculating validation accuracy, the entire network is used.

### B. Start Codon Selection

While AUG is the predominant start codon, there are several others that can be used [28]. We analyzed the predicted start codon for all nucleotide sequences and found that all sequences were found to use known start codons; namely: AUG, GUG, UUG, CUG, AUU, AUC and AUA.

### C. Predicting the Wrong Codon

Prediction of the wrong codon occurs more frequently as sequences become longer. This could be mitigated by a larger network at the cost of increased training time. The occurrence of incorrectly selecting codons can be seen in Figure 3. The axes are sorted lexicographically and, as expected, a strong band along the diagonal emerges. This is because synonymous codons often share the first two bases with the third position in the codon acting as the wobble base pair.

### D. Codons that Encode the Wrong Amino Acid

After training, the models will still choose codons that encode the wrong amino acid. It is unclear why the model sometimes selects a codon that encodes a different amino acid than what was originally provided as input. One reason may be that the network has seen so many examples in the training set of a particular motif that it has high confidence that another amino acid is more appropriate. Using the added heuristics (see Section II-E), the network is able to eliminate the selection of codons that encode the wrong amino acid and providing increased overall codon selection accuracy (see Table II).

The impact of these mis-substitutions is also unclear. Further analysis should be done to see if protein function is impacted or if changes modify the translation of this gene.

### E. Fine-Tuned Models

We fine-tuned our base model on 4 different strains of *E. coli* as well as for two other bacteria: *Staphylococcus aureus* and *Yersinia enterocolitica*. Fine-tuning was done by re-initializing the CDS and amino acid sequence classification layers and training them for 10 epochs. Results from fine-tuning can be seen in Table III.

The results of the fine-tuning show good performance on the *E. coli* bacteria and poor performance on the two other bacteria. The general many-to-one relationship of codons to amino acids was learned as evidenced by the > 99% amino acid prediction accuracy (before applying synonymous codon correction) but the codon prediction accuracy can be as low as 59.28% in the case of *Y. enterocolitica*. This suggests that the recurrent portion of our network is tied specifically to *E. coli* and does not generalize well. We had hoped that we could create a network that would generalize well but this behavior can be expected as we trained exclusively on *E. coli* data.

During fine-tuning, the network overfits to the specific strains as can be seen by the high codon prediction accuracies ( 100% for *E. coli* strains). This could be mitigated by decreasing the number of epochs that are run for fine-tuning. The appropriate number of epochs used in fine-tuning is unknown especially as the fine-tuning datasets are significantly smaller than the datasets used to train the base models (see Table III).

## IV. CONCLUSION

Here we have presented a method for representing the global codon bias of an organism. Previous methods have provided a statistical overview of codon usage and a measure of how much a particular protein deviates from a reference set of proteins. Often, these approaches are used in order to backtranslate protein sequences such that the output nucleotide sequence matches the codon bias of a reference set of proteins. Methods such as CAI, however, are too coarse and give a very high-level picture of codon usage that does not account for the local codon bias that of a protein influenced by forming structures and other phenomenon.

Our method, leveraging neural networks, language models and translation techniques used in machine-learning, provides a rich representation of codon bias by using local and long-range contexts to inform codon selection. Using this method, we can predict the correct nucleotide sequence for given proteins with high accuracy. While we have only demonstrated this ability on *E. coli*, we believe that this approach can be generalized and applied to other species allowing for more advanced methods of backtranslation.

## ACKNOWLEDGMENT

This work was funded through the Utah NASA Space Grant Consortium and EPSCoR and through the BYU Graduate Research Fellowship.

The authors would like to thank Dr. Dan Ventura, Dr. Ryan Farrell, and Dr. Amin Salehi, Nozomu Okuda, Christopher Tensemeyer, Derrall Heath, Aaron Dennis, Shayla Draper and Michael Whatkins.

## REFERENCES

- [1] P. M. Sharp and W.-H. Li, "The codon adaptation index-a measure of directional synonymous codon usage bias, and its potential applications," *Nucleic acids research*, vol. 15, no. 3, pp. 1281–1295, 1987.
- [2] H. Seligmann and D. D. Pollock, "The ambush hypothesis: hidden stop codons prevent off-frame gene reading," *DNA and cell biology*, vol. 23, no. 10, pp. 701–705, 2004.
- [3] J. R. Coleman, D. Papamichail, S. Skiena, B. Futcher, E. Wimmer, and S. Mueller, "Virus attenuation by genome-scale changes in codon pair bias," *Science*, vol. 320, no. 5884, pp. 1784–1787, 2008.
- [4] S. L. Chen, W. Lee, A. K. Hottes, L. Shapiro, and H. H. McAdams, "Codon usage between genomes is constrained by genome-wide mutational processes," *Proceedings of the National Academy of Sciences of the United States of America*, vol. 101, no. 10, pp. 3480–3485, 2004.
- [5] P. M. Sharp, L. R. Emery, and K. Zeng, "Forces that influence the evolution of codon bias," *Philosophical Transactions of the Royal Society B: Biological Sciences*, vol. 365, no. 1544, pp. 1203–1212, 2010.
- [6] A. Hilterbrand, J. Saelens, and C. Putonti, "Cbdb: the codon bias database," *BMC bioinformatics*, vol. 13, no. 1, p. 62, 2012.
- [7] H. M. Salim and A. R. Cavalcanti, "Factors influencing codon usage bias in genomes," *Journal of the Brazilian Chemical Society*, vol. 19, no. 2, pp. 257–262, 2008.
- [8] E. M. Novoa and L. R. de Pouplana, "Speeding with control: codon usage, tRNAs, and ribosomes," *Trends in Genetics*, vol. 28, no. 11, pp. 574–581, 2012.
- [9] J. B. Plotkin and G. Kudla, "Synonymous but not the same: the causes and consequences of codon bias," *Nature Reviews Genetics*, vol. 12, no. 1, pp. 32–42, 2011.
- [10] P. S. Spencer and J. M. Barral, "Genetic code redundancy and its influence on the encoded polypeptides," *Computational and structural biotechnology journal*, vol. 1, no. 1, pp. 1–8, 2012.

TABLE III. FINE-TUNING RESULTS USING THE SEQUENCE LENGTH 500, NO DROPOUT MODEL WITH ALL WEIGHTS FROZEN EXCEPT FOR THE PREDICTION LAYERS. PREDICTION LAYER WEIGHTS ARE NOT RESET. REPORTED ACCURACIES ARE FOR CODON SELECTION ACCURACIES ONLY. ALL AMINO ACID ACCURACIES WERE > 99%. THE DATASETS FOR FINE-TUNING ARE DRAMATICALLY SMALLER THAN THOSE USED TO TRAIN THE BASE MODEL. THE NUMBER OF SEQUENCES FOR EACH DATASET IS REPORTED, 80% OF THE DATASET WAS USED FOR TRAINING/VALIDATION AND THE OTHER 20% USED FOR TESTING.

Name	# seqs	Training	Validation	Testing
<i>E. coli</i> str. K-12 substr. MG1655 (NC_000913.3)	3570	100.00%	95.64%	95.73%
<i>E. coli</i> UTI89 (NC_007946.1)	4243	99.98%	91.66%	91.99%
<i>E. coli</i> 536 (NC_008253.1)	4012	99.98%	90.18%	89.43%
<i>E. coli</i> APEC O1 (NC_008563.1)	4307	99.98%	91.91%	91.59%
<i>Y. enterocolitica</i> subsp. <i>enterocolitica</i> 8081 (NC_008800.1)	3385	87.74%	45.22%	46.78%
<i>S. aureus</i> RF122 (NC_007622.1)	2318	81.19%	59.09%	59.28%

- [11] M. Welch, S. Govindarajan, J. E. Ness, A. Villalobos, A. Gurney, J. Minshull, and C. Gustafsson, "Design parameters to control synthetic gene expression in escherichia coli," *PloS one*, vol. 4, no. 9, p. e7002, 2009.
- [12] R. Hershberg and D. A. Petrov, "Selection on codon bias," *Annual review of genetics*, vol. 42, pp. 287–299, 2008.
- [13] T. Ikemura, "Codon usage and tRNA content in unicellular and multicellular organisms..," *Molecular biology and evolution*, vol. 2, no. 1, pp. 13–34, 1985.
- [14] F. Yamao, Y. Andachi, A. Muto, T. Ikemura, and S. Osawa, "Levels of trnas in bacterial cells as affected by amino acid usage in proteins," *Nucleic acids research*, vol. 19, no. 22, pp. 6119–6122, 1991.
- [15] M. Allert, J. C. Cox, and H. W. Hellinga, "Multifactorial determinants of protein expression in prokaryotic open reading frames," *Journal of molecular biology*, vol. 402, no. 5, pp. 905–918, 2010.
- [16] S. Pechmann and J. Frydman, "Evolutionary conservation of codon optimality reveals hidden signatures of cotranslational folding," *Nature structural & molecular biology*, vol. 20, no. 2, pp. 237–243, 2013.
- [17] D. B. Goodman, G. M. Church, and S. Kosuri, "Causes and effects of n-terminal codon bias in bacterial genes," *Science*, vol. 342, no. 6157, pp. 475–479, 2013.
- [18] K. Bentele, P. Saffert, R. Rauscher, Z. Ignatova, and N. Blüthgen, "Efficient translation initiation dictates codon usage at gene start," *Molecular systems biology*, vol. 9, no. 1, p. 675, 2013.
- [19] D. M. Hoover and J. Lubkowski, "Dnaworks: an automated method for designing oligonucleotides for pcr-based gene synthesis," *Nucleic acids research*, vol. 30, no. 10, pp. e43–e43, 2002.
- [20] M. Schuster and K. K. Paliwal, "Bidirectional recurrent neural networks," *IEEE Transactions on Signal Processing*, vol. 45, no. 11, pp. 2673–2681, 1997.
- [21] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [22] J. Chung, C. Gulcehre, K. Cho, and Y. Bengio, "Empirical evaluation of gated recurrent neural networks on sequence modeling," *arXiv preprint arXiv:1412.3555*, 2014.
- [23] P. J. Cock, T. Antao, J. T. Chang, B. A. Chapman, C. J. Cox, A. Dalke, I. Friedberg, T. Hamelryck, F. Kauff, B. Wilczynski, *et al.*, "Biopython: freely available python tools for computational molecular biology and bioinformatics," *Bioinformatics*, vol. 25, no. 11, pp. 1422–1423, 2009.
- [24] I. Sutskever, O. Vinyals, and Q. V. Le, "Sequence to sequence learning with neural networks," in *Advances in neural information processing systems*, pp. 3104–3112, 2014.
- [25] D. Bahdanau, K. Cho, and Y. Bengio, "Neural machine translation by jointly learning to align and translate," *arXiv preprint arXiv:1409.0473*, 2014.
- [26] J. Donahue, Y. Jia, O. Vinyals, J. Hoffman, N. Zhang, E. Tzeng, and T. Darrell, "Decaf: A deep convolutional activation feature for generic visual recognition," *CoRR*, vol. abs/1310.1531, 2013.
- [27] A. V. Lobanov, A. A. Turanov, D. L. Hatfield, and V. N. Gladyshev, "Dual functions of codons in the genetic code," *Critical Reviews in Biochemistry and Molecular Biology*, vol. 45, no. 4, pp. 257–265, 2010.
- [28] A. Hecht, J. Glasgow, P. R. Jaschke, L. Bawazer, M. S. Munson, J. Cochran, D. Endy, and M. Salit, "Measurements of translation initiation from all 64 codons in *e. coli*," *bioRxiv*, p. 063800, 2016.
- [29] F. Chollet, "Keras." <https://github.com/fchollet/keras>, 2015.
- [30] R. Al-Rfou, G. Alain, A. Almhairi, C. Angermueller, D. Bahdanau, N. Ballas, F. Bastien, J. Bayer, A. Belikov, A. Belopolsky, Y. Bengio, A. Bergeron, J. Bergstra, V. Bisson, J. Blecher-Snyder, N. Bouchard, N. Boulanger-Lewandowski, X. Bouthillier, A. de Brébisson, O. Breuleux, P.-L. Carrier, K. Cho, J. Chorowski, P. Christiano, T. Cooijmans, M.-A. Côté, A. Courville, Y. N. Dauphin, O. Delalleau, J. Demouth, G. Desjardins, S. Dieleman, L. Dinh, M. Ducoffe, V. Dumoulin, S. Ebrahimi Kahou, D. Erhan, Z. Fan, O. Firat, M. Germain, X. Glorot, I. Goodfellow, M. Graham, C. Gulcehre, P. Hamel, I. Harlouchet, J.-P. Heng, B. Hidasi, S. Honari, A. Jain, S. Jean, K. Jia, M. Korobov, V. Kulkarni, A. Lamb, P. Lamblin, E. Larsen, C. Laurent, S. Lee, S. Lefrancois, S. Lemieux, N. Léonard, Z. Lin, J. A. Livezey, C. Lorenz, J. Lowin, Q. Ma, P.-A. Manzagol, O. Mastropietro, R. T. McGibbon, R. Memisevic, B. van Merriënboer, V. Michalski, M. Mirza, A. Orlandi, C. Pal, R. Pascanu, M. Pezeshki, C. Raffel, D. Renshaw, M. Rocklin, A. Romero, M. Roth, P. Sadowski, J. Salvatier, F. Savard, J. Schlüter, J. Schulman, G. Schwartz, I. V. Serban, D. Serdyuk, S. Shabanian, E. Simon, S. Spieckermann, S. R. Subramanyam, J. Sygnowski, J. Tanguay, G. van Tulder, J. Turian, S. Urban, P. Vincent, F. Visin, H. de Vries, D. Warde-Farley, D. J. Webb, M. Willson, K. Xu, L. Xue, L. Yao, S. Zhang, and Y. Zhang, "Theano: A Python framework for fast computation of mathematical expressions," *arXiv e-prints*, vol. abs/1605.02688, May 2016.
- [31] D. Kingma and J. Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.